

# Responsible AI: Deep Learning Model Verification

## Progress Review - 3

Candidate: Guanqin Zhang  
Main Supervisor: Yulei Sui,  
Co-Supervisor: Dilum Bandara, Shiping Chen

University of New South Wales, Sydney, Australia  
CSIRO, Data61

# Review Agenda

- ① **Part I:** Introduction of Neural Network Verification
- ② **Part II:** Progress Review Summary
- ③ **Part III:** Main Contribution of Works During PhD Candidature (leading as **first-authored** paper)
  - A Tale of Two Cities: Data and Configuration Variances in Robust Deep Learning (IEEE Computer Society)
  - Efficient Neural Network Verification via Order Leading Exploration of Branch-and-Bound Trees (ECOOP 2025) - CORE A
  - Adaptive Branch-and-Bound Tree Exploration for Neural Network Verification (DATE 2025) - CORE B
  - Efficient Incremental Verification of Neural Networks Guided by Counterexample Potentially (OOPSLA 2025) - CORE A
- ④ **Part IV:** Conclusion

# Introduction Neural Network Verification

## Progress Review - 3: Part I

Candidate: Guanqin Zhang  
Main Supervisor: Yulei Sui,  
Co-Supervisor: Dilum Bandara, Shiping Chen

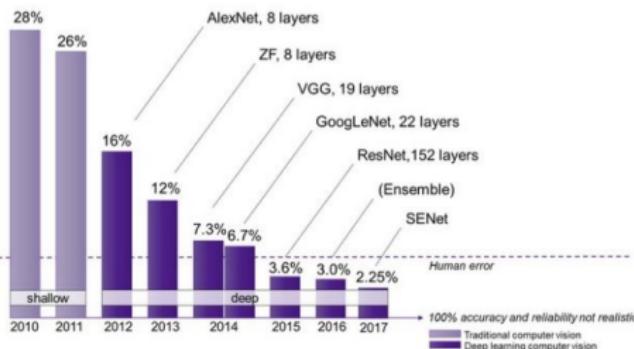
University of New South Wales, Sydney, Australia  
CSIRO, Data61

# Part I: Introduction of Neural Network Verification

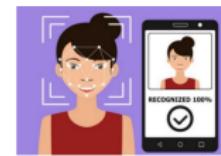
- ① Background
- ② Local Robustness
- ③ Verification in Practice
- ④ Branch & Bound

# Modern AI Models are Large and Complex

Significant Advancement of Deep Neural Networks



Autonomous Driving Vehicles



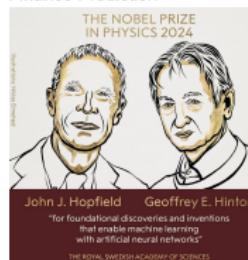
Facial Recognition of Mobile Payment System



Finance Prediction



Speech Recognition



John J. Hopfield      Geoffrey E. Hinton  
"for foundational discoveries and inventions that enable machine learning with artificial neural networks"



David Baker      Demis Hassabis      John M. Jumper  
"for computational protein design"      "for protein structure prediction"

# Robustness Verification is Needed in all Hands

**Correctness Issue:** Misclassification

Original image: Dermatoscopic image of a benign melanocytic nevus, along with the diagnostic probability computed by a deep neural network.  
Adversarial noise: Perturbation computed by a common adversarial attack technique. See (7) for details.  
Adversarial example: Combined image of nevus and attack perturbation and the diagnostic probabilities from the same deep neural network.

Model confidence: Benign Malignant

**Security Issue:** DeepFake

real fake  
real fake

**Robustness Issue:** fragile to the imperceptible noise

original time series  $X$ : class-1 with 99% confidence  
imperceptible noise:  $\eta = \epsilon \cdot \text{sign}(\nabla_x J(X, \hat{Y}))$   
perturbed time series  $X + \eta$ : class-2 with 99% confidence

**Privacy Issue:** Imprecise facial recognition

Perturbing frames to impersonate

**Safety Issue:** incorrect recognition

Uniform illumination: Stop Sign  
Our illumination: Stop Sign  
Captured: Stop Sign  
Speed 30

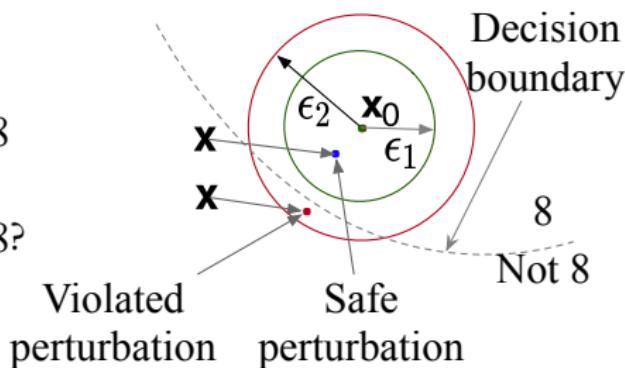
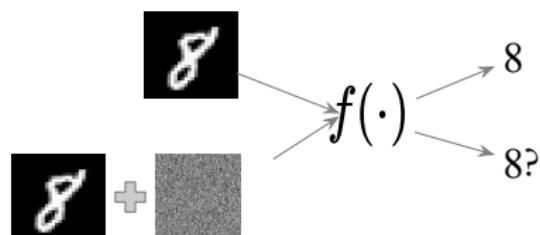
Impersonating Mila Jovovich (by Georges Biard)

**Small perturbations** in the input space can **fool** the neural networks to yield **incorrect output**.

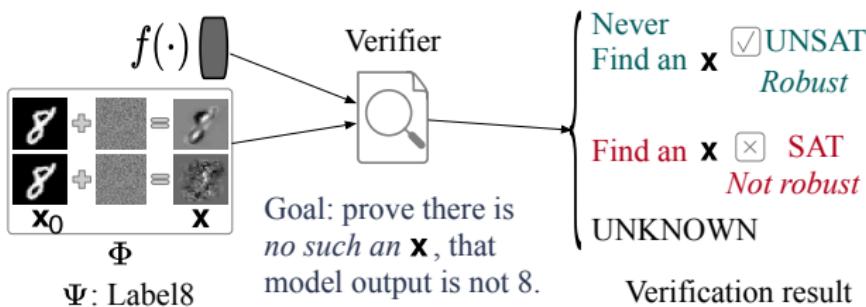
# Verification for Local Robustness

Given a model  $f(\cdot)$ , an input  $\mathbf{x}_0$ , and a perturbation radius  $\epsilon$ , the model is robust for all perturbed  $\mathbf{x}$ :

Robustness

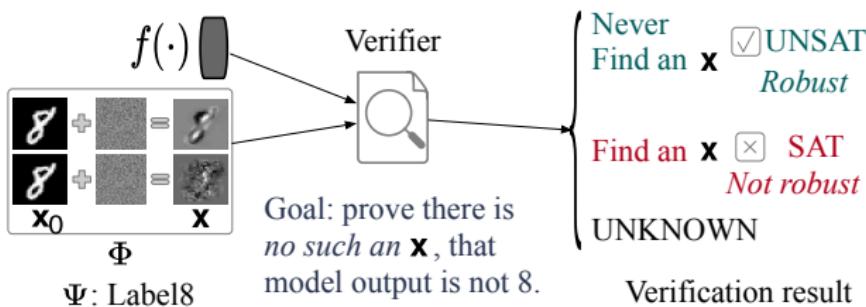


# Verification Procedure



- Verification aims to certify “specification guarantees” on model behaviours by rigorously proving the absence of any specification violations.

# Verification Procedure



- Verification aims to certify “specification guarantees” on model behaviours by rigorously proving the absence of any specification violations.
- Given an input  $\mathbf{x} \in \mathcal{I}$  of a network  $f(\cdot)$  to be verified, and a specification consisting of an input property  $\Phi$  and an output property  $\Psi$ , the verification is to check  $\mathbf{x} \models \Phi \implies f(\mathbf{x}) \models \Psi$ .

# Neural Network Testing v.s. Verification

**Testing** evaluates DNN behaviour, while **verification** proves formal guarantees

- **Methodology:** Testing uses **empirical** methods, while verification employs **formal** techniques.

# Neural Network Testing v.s. Verification

**Testing** evaluates DNN behaviour, while **verification** proves formal guarantees

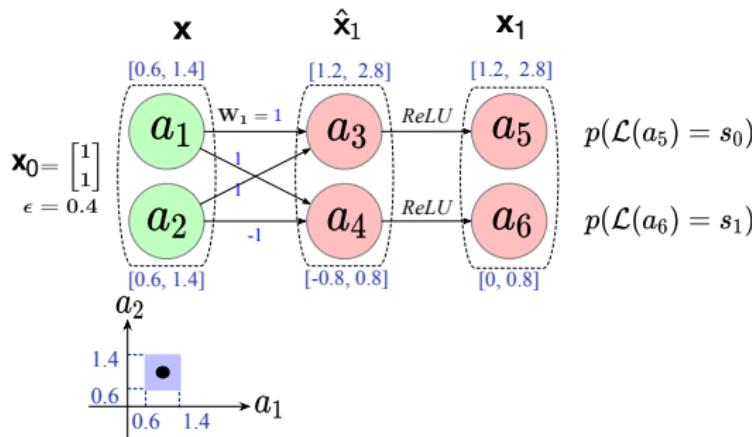
- **Methodology:** Testing uses **empirical** methods, while verification employs **formal** techniques.
- **Completeness:** Testing uncovers issues via **case by case**, but doesn't **guarantee** the absence of errors, whereas verification aims to prove correctness.

# Neural Network Testing v.s. Verification

**Testing** evaluates DNN behaviour, while **verification** proves formal guarantees

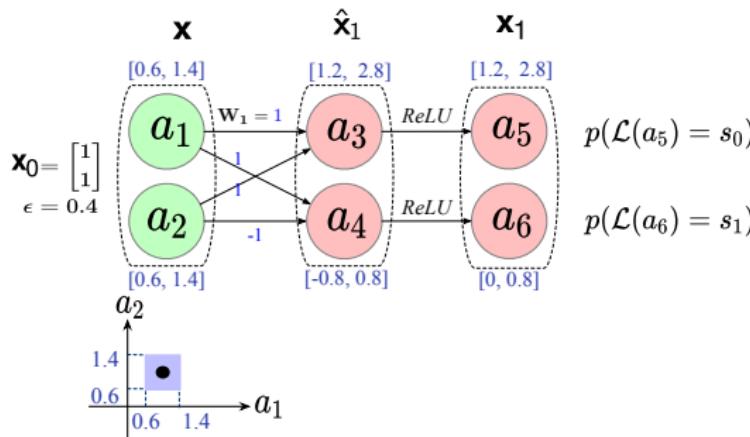
- **Methodology:** Testing uses **empirical** methods, while verification employs **formal** techniques.
- **Completeness:** Testing uncovers issues via **case by case**, but doesn't **guarantee** the absence of errors, whereas verification aims to prove correctness.
- **Scalability:** Testing is more scalable, while verification can be computationally intensive since it aims to approximate **all** possible scenarios to provide a correctness guarantee.

# Verification in Practice



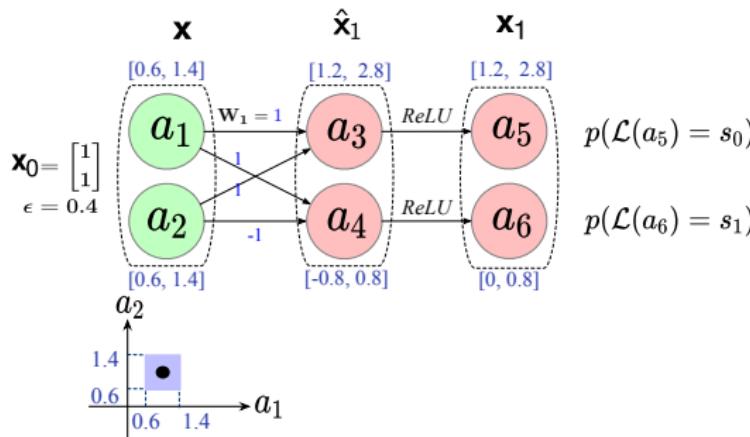
- Perturbed space is  $C_{in} = 0.6 \leq a_1 \leq 1.4 \wedge 0.6 \leq a_2 \leq 1.4$

# Verification in Practice



- Perturbed space is  $C_{in} = 0.6 \leq a_1 \leq 1.4 \wedge 0.6 \leq a_2 \leq 1.4$
- Expected output condition is  $C_{out} = a_5 > a_6$

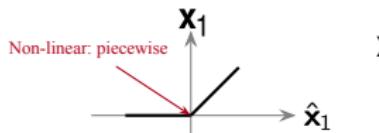
# Verification in Practice



- Perturbed space is  $C_{in} = 0.6 \leq a_1 \leq 1.4 \wedge 0.6 \leq a_2 \leq 1.4$
- Expected output condition is  $C_{out} = a_5 > a_6$
- The neural network is  $C_f = \bigwedge_{i=1}^1 ReLU(\sum_{j=1}^1 \mathbf{Wx} + b)$ . (In our example,  $b = 0$ )

# Why Neural Network Verification is Hard?

*ReLU* functions is **piece-wise**, i.e.,  $\text{ReLU}(x) = \begin{cases} x, & x > 0 \\ 0, & x \leq 0 \end{cases}$



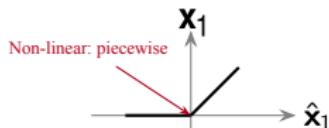
$$\mathbf{x}_1 \equiv ((\hat{\mathbf{x}}_1 \leq 0) \wedge (\mathbf{x}_1 = 0)) \vee (\hat{\mathbf{x}}_1 > 0) \wedge (\mathbf{x}_1 = \hat{\mathbf{x}}_1),$$

**bisects** in the verification problem.

<sup>6</sup>Katz *et al.* Reluplex: An efficient SMT solver for verifying deep neural networks. CAV 2017

# Why Neural Network Verification is Hard?

*ReLU* functions is **piece-wise**, i.e.,  $\text{ReLU}(x) = \begin{cases} x, & x > 0 \\ 0, & x \leq 0 \end{cases}$



$$\mathbf{x}_1 \equiv ((\hat{\mathbf{x}}_1 \leq 0) \wedge (\mathbf{x}_1 = 0)) \vee (\hat{\mathbf{x}}_1 > 0) \wedge (\mathbf{x}_1 = \hat{\mathbf{x}}_1),$$

**bisects** in the verification problem.

These **disjunction** constraint  $C_{in} \wedge C_f \wedge \neg C_{out} =$

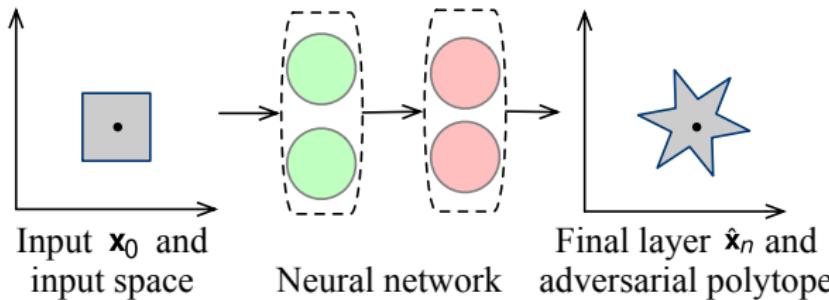
$$\begin{aligned}
 0.6 \leq a_1 \leq 1.4 \wedge 0.6 \leq a_2 \leq 1.4 \wedge & (a_1 + a_2 = a_3) \wedge (a_1 - a_2 = a_4) \wedge \\
 ((a_3 \leq 0) \wedge (a_5 = 0)) \vee & (a_3 > 0) \wedge (a_5 = a_3) \wedge \\
 (a_4 \leq 0) \wedge (a_6 = 0) \vee & (a_4 > 0) \wedge (a_6 = a_4) \wedge \\
 a_5 \leq a_6
 \end{aligned} \tag{1}$$

The SAT solving spends **exponential** time to find all variables assignments that satisfy all disjunct clauses<sup>6</sup>. It is an **NP-hard** problem.

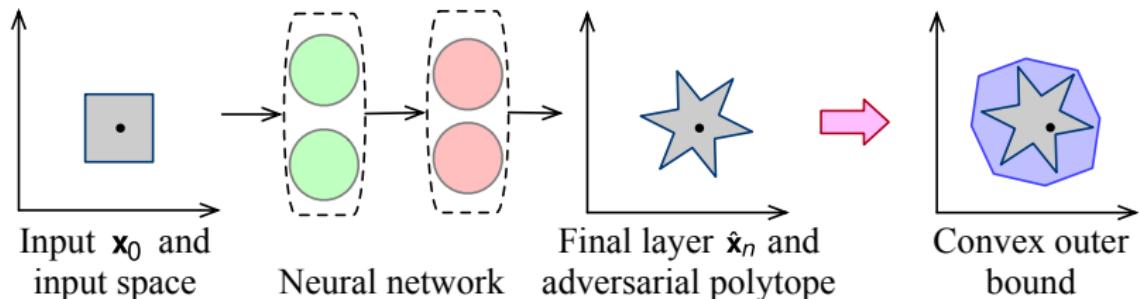
<sup>6</sup> Katz et al. Reluplex: An efficient SMT solver for verifying deep neural networks. CAV 2017

# Why Neural Network Verification is Hard?

- With **higher** dimensions and **deeper** layers, the **composed** nonlinear activation functions (*ReLU, Sigmod, tanh*) , operate on **non-convex polytope** hyperspace (**non-convex** NN function  $f(\cdot)$ ).
- The verification process is too computationally costly for non-convex polytope. It needs to approximate tractable bounds to form a provable convex solution space.



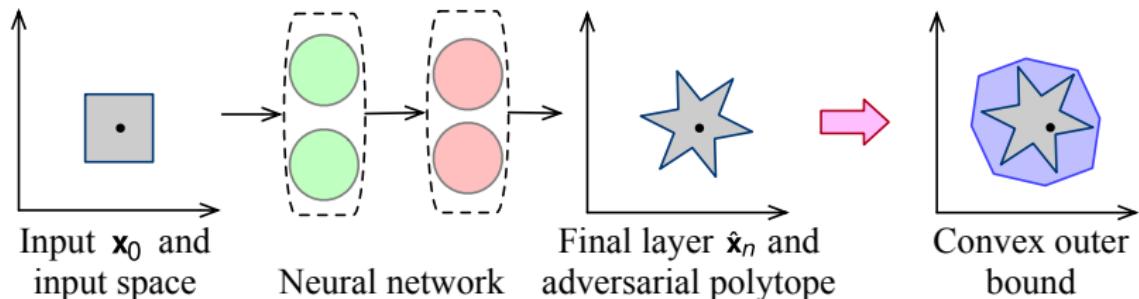
# Constraint Extraction: Convex Relaxation



- **Adversarial polytope** is the set of all final-layer activations that can be achieved by applying a norm-bounded perturbation to the input, which is **non-convex**;

<sup>1</sup>Liu, Changliu, et al. "Algorithms for verifying deep neural networks." Foundations and Trends in Optimization 2021

# Constraint Extraction: Convex Relaxation

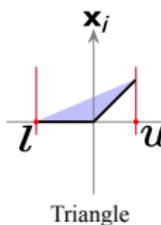
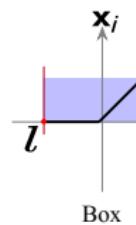
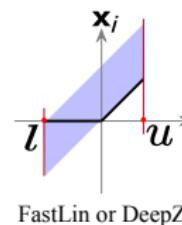
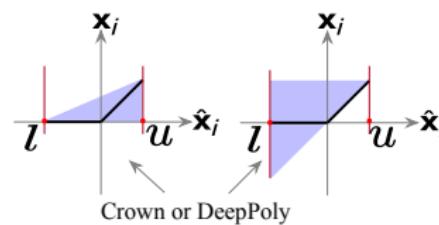


- **Adversarial polytope** is the set of all final-layer activations that can be achieved by applying a norm-bounded perturbation to the input, which is **non-convex**;
- **Convex outer bound** is an over-approximation of the adversarial polytope with linear convexity<sup>1</sup>. Robustness is guaranteed if the classification results do not change (unaffected by perturbations) within the **convex outer bound**

<sup>1</sup>Liu, Changliu, et al. "Algorithms for verifying deep neural networks." Foundations and Trends in Optimization 2021

# Linear Relaxation (*ReLU*)

Approaches to approximating the *ReLU* to **linear convex** shapes (the blue area)<sup>2</sup>:

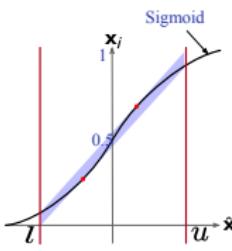
Ehlers *et al.* ATVA 2017Gehr *et al.* S&P 2018Weng *et al.* ICML 2018Crown: Zhang *et al.* NeurIPS 2018   DeepPoly: Singh *et al.* POPL 2019

<sup>2</sup>Triangle: Ehlers *et al.*; Box: Gehr *et al.*; Zonotope: Fast-Lin and Fast-Lip: Weng *et al.*; DeepPoly: Singh *et al.*; Crown: Zhang *et al.*



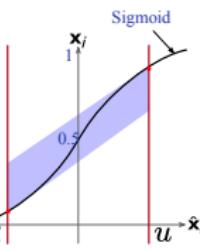
# Relaxation (Sigmoid)

- Differential functions, such as  $\text{Sigmoid}(\mathbf{x}) = \frac{1}{e^{\mathbf{x}} + 1}$ , are '**S-shaped**' curves, which is convex for values less than a **particular point** (e.g., 0), and concave for values greater than that point.
- They can be approximated to a convex region when we decide on values less than a **particular point**, and bounded with different strategies:



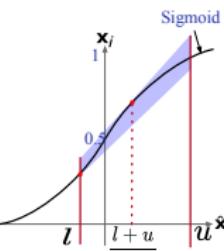
Minimal area and parallel lines

Henriksen et al. ECAI 2020



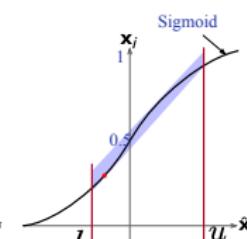
Endpoints

Zhang et al. ASE 2022



Minimal area

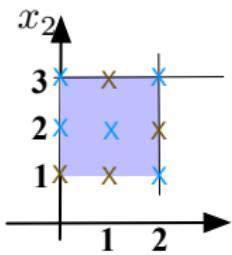
Henriksen et al. ECAI 2020



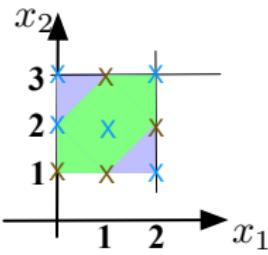
Parallel lines

Wu et al. AAAI 2022

# Interval & Zonotop I



**Interval**



**Zonotop**

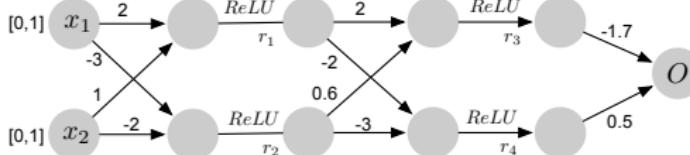
- Brown points are from the concrete domain.
- Blue points are the approximated points from the abstract domain.

- $x_1, x_2 \in \mathbb{Z}$
- Interval space takes 9 integer points, while Zonotope takes 7 points.

# A Case Study by BaB-based Sound and Complete DNN Verification

Specification:  $\Phi \wedge \Psi$     Input:  $\Phi = x_1 \in [0, 1] \wedge x_2 \in [0, 1]$     Output:  $\Psi = (O \geq 0)$

Network:  $N$



To calculate:  
 $\text{Minimize}(O) \rightarrow \text{LP}(\Phi \wedge \Psi \wedge N)$

Sound but incomplete

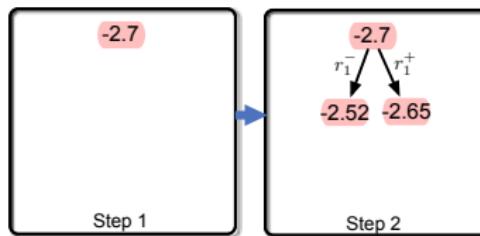
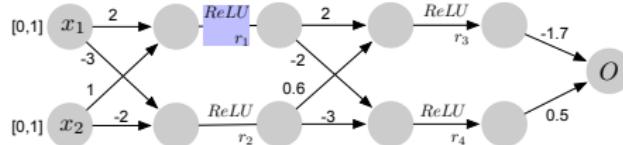
-2.7

Step 1

# A Case Study by BaB-based Sound and Complete DNN Verification

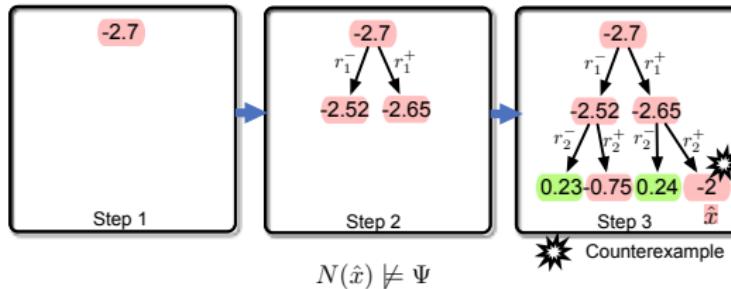
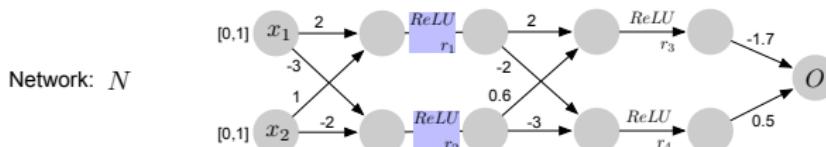
Specification:  $\Phi \wedge \Psi$     Input:  $\Phi = x_1 \in [0, 1] \wedge x_2 \in [0, 1]$     Output:  $\Psi = (O \geq 0)$

Network:  $N$



# A Case Study by BaB-based Sound and Complete DNN Verification

Specification:  $\Phi \wedge \Psi$     Input:  $\Phi = x_1 \in [0, 1] \wedge x_2 \in [0, 1]$     Output:  $\Psi = (O \geq 0)$



Specification Tree generated after Branch and Bound (BaB), when verifier cannot certify the network.

# Progress Review Summary

## Progress Review - 3: Part II

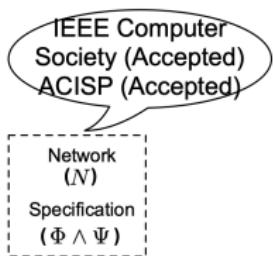
Candidate: Guanqin Zhang

Main Supervisor: Yulei Sui,

Co-Supervisor: Dilum Bandara, Shiping Chen

University of New South Wales, Sydney, Australia  
CSIRO, Data61

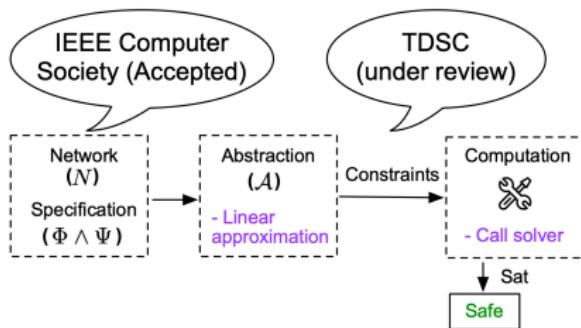
# A High-level View of Neural Network Verification Procedures<sup>3 4</sup>



<sup>3</sup>IEEE Computer: A tale of two cities: data and configuration variances in robust deep learning, **Guanqin Zhang**, J Sun, F Xu, Y Sui, HMND Bandara, S Chen, T Menzies

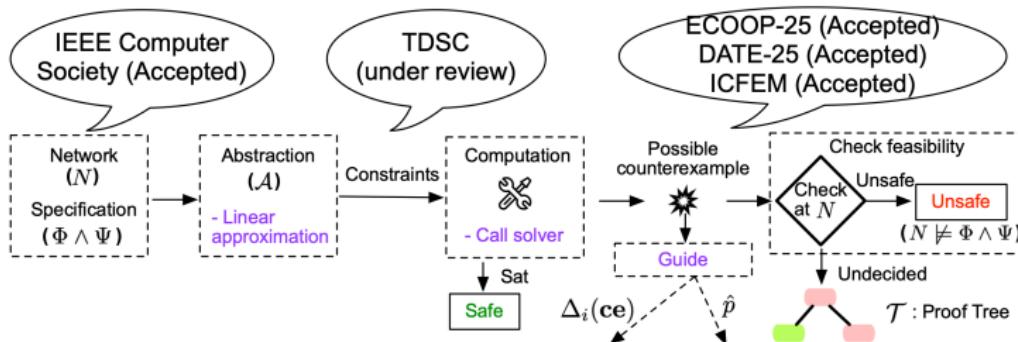
<sup>4</sup>ACISP: Understanding the Robustness of Machine Unlearning Models: **Guanqin Zhang\***, F Xu, Dilum Bandara, Shiping Chen and Yulei Sui (conditionally Accept)

# A High-level View of Neural Network Verification Procedures<sup>5</sup>



<sup>5</sup>Unveiling the Boundaries of Neural Network Verifiers: A Study on Soundness and Completeness, Feng Xu\*, **Guanqin Zhang\***, Zhenya Zhang, Xun Li, Ivor W Tsang, Yew-Soon Ong, and Yulei Sui

# A High-level View of Neural Network Verification Procedures<sup>6 7 8</sup>



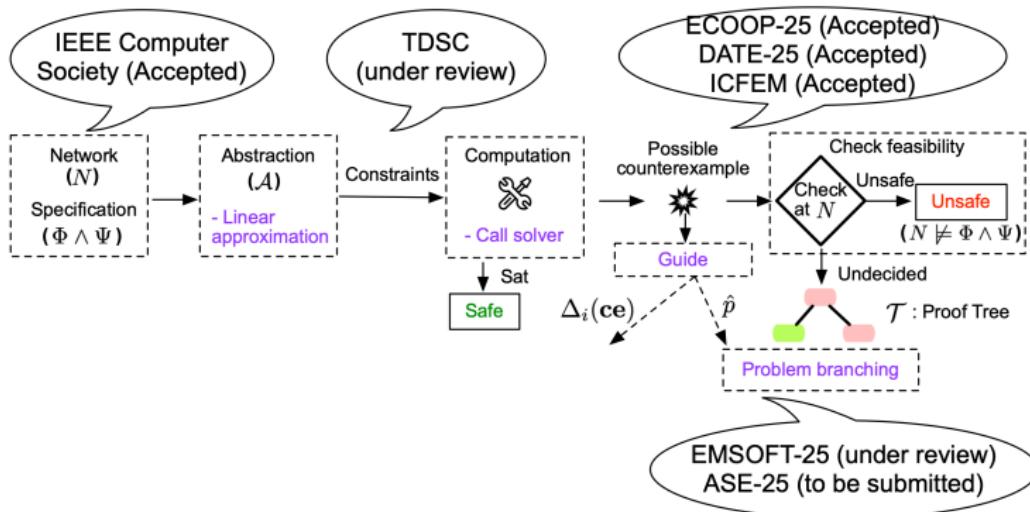
<sup>6</sup> ECOOP25: Efficient Neural Network Verification via Order Leading Exploration of Branch-and-Bound Trees : **Guanqin Zhang**, Kota Fukuda, Zhenya Zhang, Dilum Bandara, Shiping Chen, Jianjun Zhao, Yulei Sui

<sup>7</sup> DATE25: Adaptive Branch-and-Bound Tree Exploration for Neural Network Verification: Kota Fukuda, **Guanqin Zhang**, Zhenya Zhang, Yulei Sui, Jianjun Zhao

<sup>8</sup> ICFM25: Eager to Stop: Efficient Falsification of Deep Neural Networks: **Guanqin Zhang**

# A High-level View of Neural Network Verification Procedures

<sup>9 10</sup>

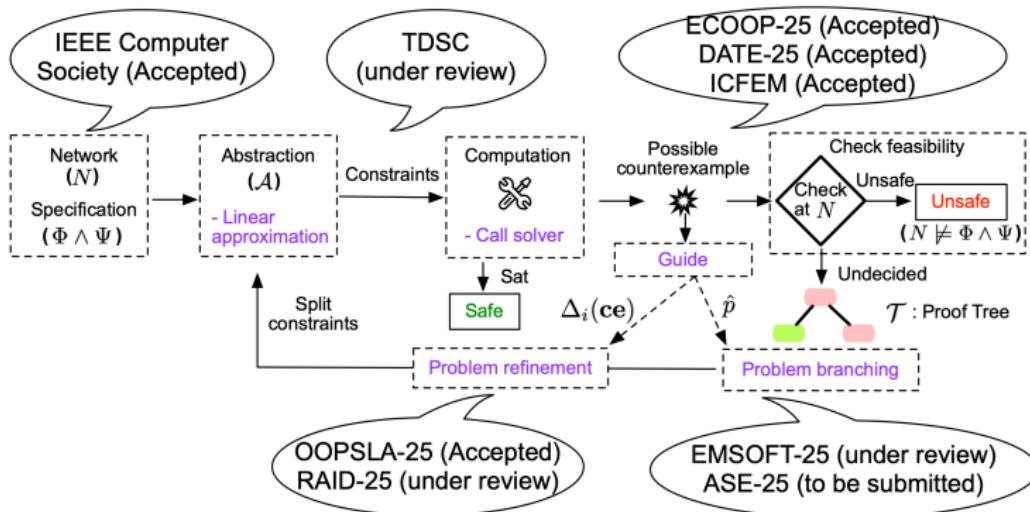


<sup>9</sup>EMSOFT25: Spurious Counterexample-Guided Branch and Bound Refinement for Neural Network Verification: **Guanqin ZHANG\***, Jiawei Ren\*, Zhenya Zhang, Yulei Sui

<sup>10</sup>ASE25: Efficient Neural Network Verification by Verdict Boundary Mining: Jiawei Ren\*, **Guanqin ZHANG\***, Zhenya Zhang, Yulei Sui

# A High-level View of Neural Network Verification Procedures

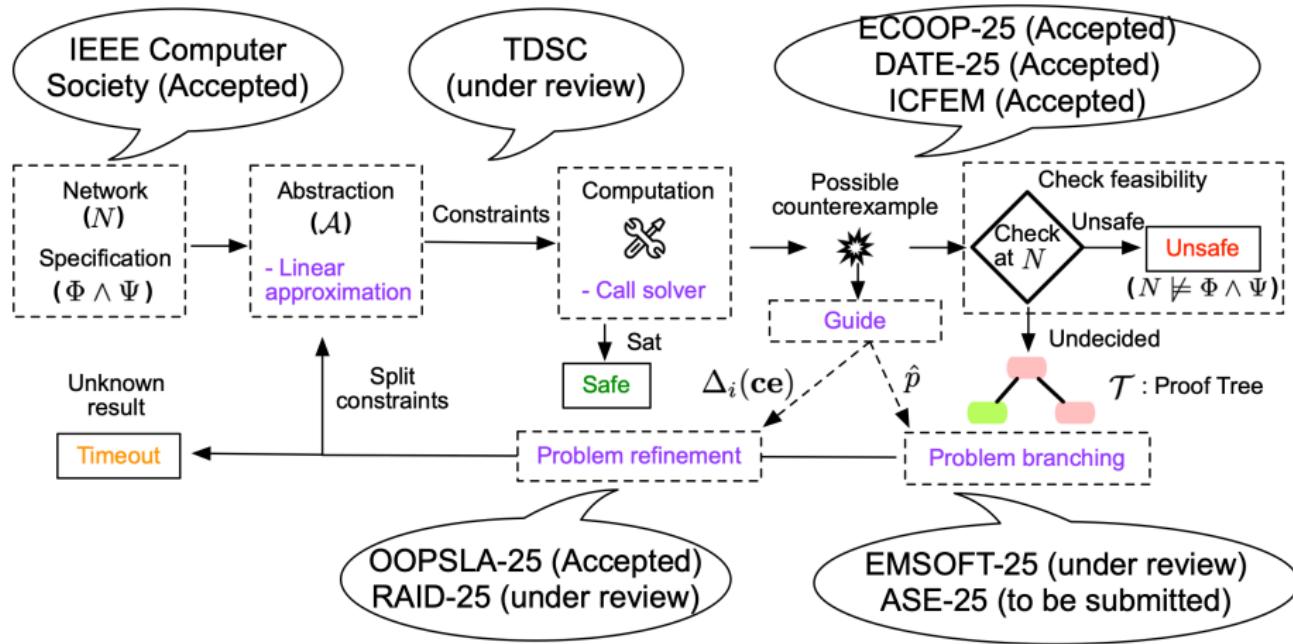
<sup>11</sup> <sup>12</sup>



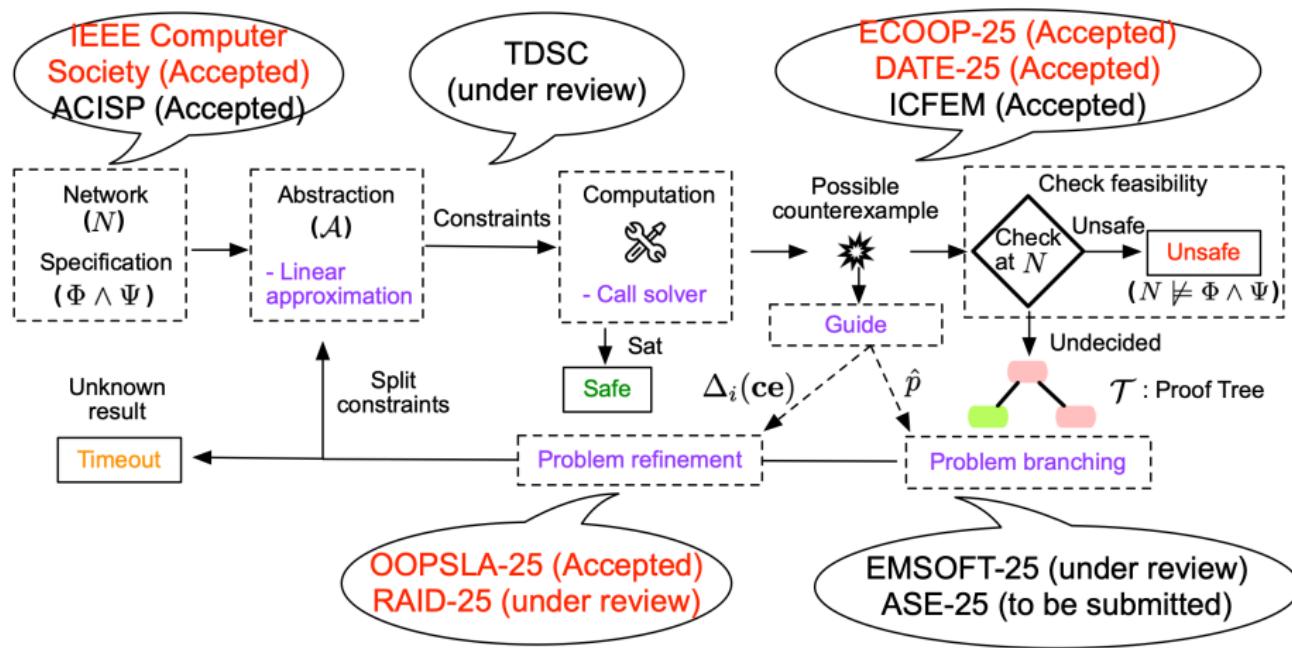
<sup>11</sup> OOSPLA25: Efficient Incremental Verification of Neural Networks Guided by Counterexample Potentially: **Guanqin Zhang**, Zhenya Zhang, Dilum Bandara, Shiping Chen, Jianjun Zhao, Yulei Sui

<sup>12</sup> RAID25: Adaptive Branch and Bound Scheduling for Incremental Neural Network

# A High-level View of Neural Network Verification Procedures



# A High-level View of Neural Network Verification Procedures



# Thesis of Neural Network Verification

## Progress Review - 3: Part III

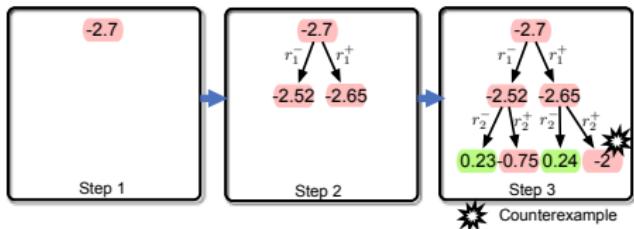
Candidate: Guanqin Zhang  
Main Supervisor: Yulei Sui,  
Co-Supervisor: Dilum Bandara, Shiping Chen

University of New South Wales, Sydney, Australia  
CSIRO, Data61

# Table of Contents

- ⑥ Greedy: Fast Ordered Exploration
  - ⑦ RL: Adaptive Verification
  - ⑧ ML: Stochastic Approach
  - ⑨ Incremental Verification
  - ⑩ PhD Justification

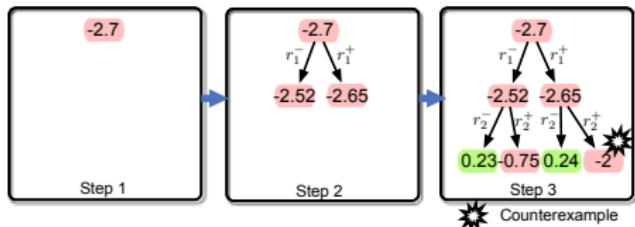
Order Leading Verification Approach (Oliva)<sup>6</sup>



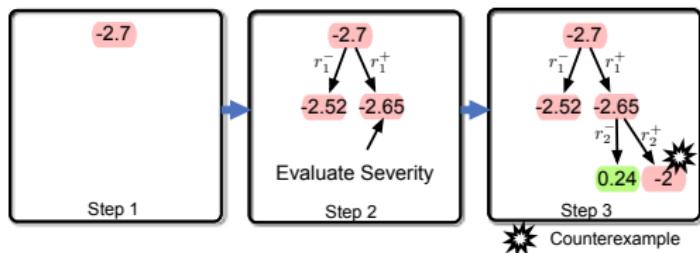
The conventional BaB algorithm manages its subproblems using a **queue** data structure. This approach results in a "first-come, first-served" processing order .

<sup>6</sup>ECOOP 2025: Guanqin ZHANG(University of New South Wales), Kota Fukuda(Kyushu University), Zhenya Zhang(Kyushu University), Dilum Bandara(Data61 CSIRO), Shiping Chen(Data61 CSIRO), Yulei Sui(University of New South Wales), Jianjun Zhao(Kyushu University)

Order Leading Verification Approach (Oliva)<sup>6</sup>



The conventional BaB algorithm manages its subproblems using a **queue** data structure. This approach results in a "first-come, first-served" processing order .



OLIVA considers the **severity** of the sub-problems, such that speedup the verification process and save **1** times of branching, and **2** times of bounding.

<sup>6</sup>ECOOP 2025: Guanqin ZHANG(University of New South Wales), Kota Fukuda(Kyushu University), Zhenya Zhang(Kyushu University), Dilum Bandara(Data61 CSIRO), Shiping Chen(Data61 CSIRO), Yulei Su(University of New South Wales), Jianjun Zhao(Kyushu University)

# Counterexample Potentially (CEPO)

- specification distance  $\hat{p}$

→ The nodes with smaller  $\hat{p}$  are likely to have high CeP.

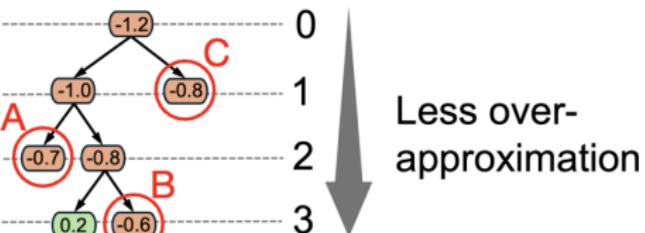
  - node depth:

over-approximation level compared to other nodes at different depths.

  - Definition of CeP

0  
1  
2  
3

Less over-approximation



$$[\![\Gamma]\!] := \begin{cases} -\infty & \text{if } \hat{p} > 0 \\ +\infty & \text{if } \hat{p} < 0 \text{ and valid}(\hat{x}) \\ \lambda \frac{\text{depth}(\Gamma)}{K} + (1 - \lambda) \frac{\hat{p}}{\hat{p}_{min}} & \text{otherwise} \end{cases}$$

- A: -0.62
- B: -0.65
- C: -0.64

B is  
selected

## Monte Carlo tree search (MCTS)

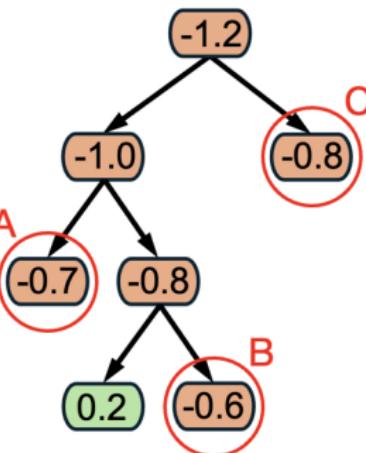
- Monte Carlo Tree Search (MCTS) effectively guides node expansion while addressing scenarios where selection based on CeP alone is difficult.

### Condition:

- Node expansion options: A, B, and C have close CeP values.
  - The left side: exploited and likely contains counterexamples.
  - The right side: unexplored but may also contain counterexamples.

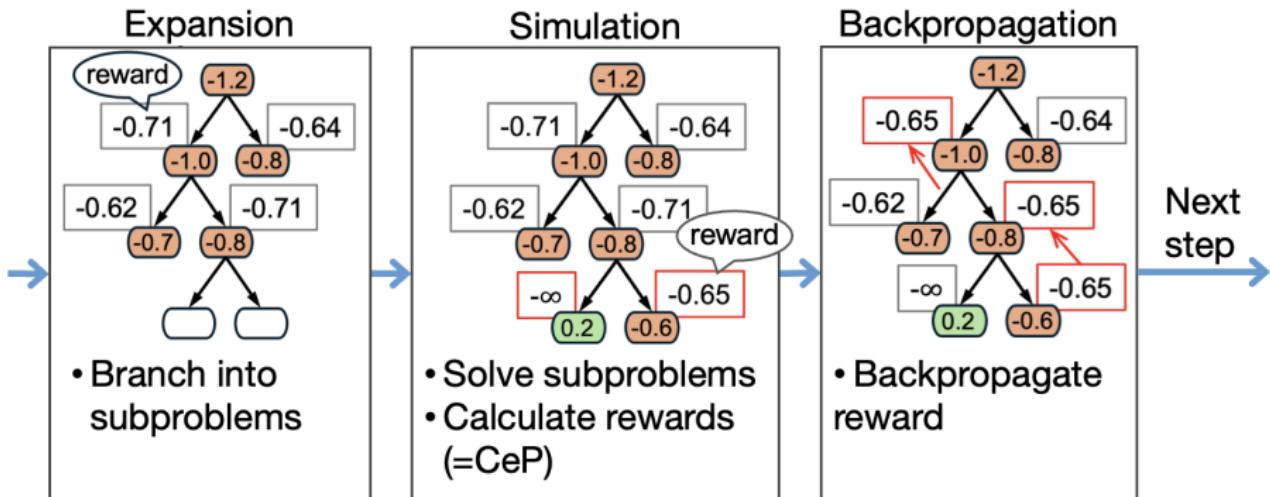


## What is the best choice?

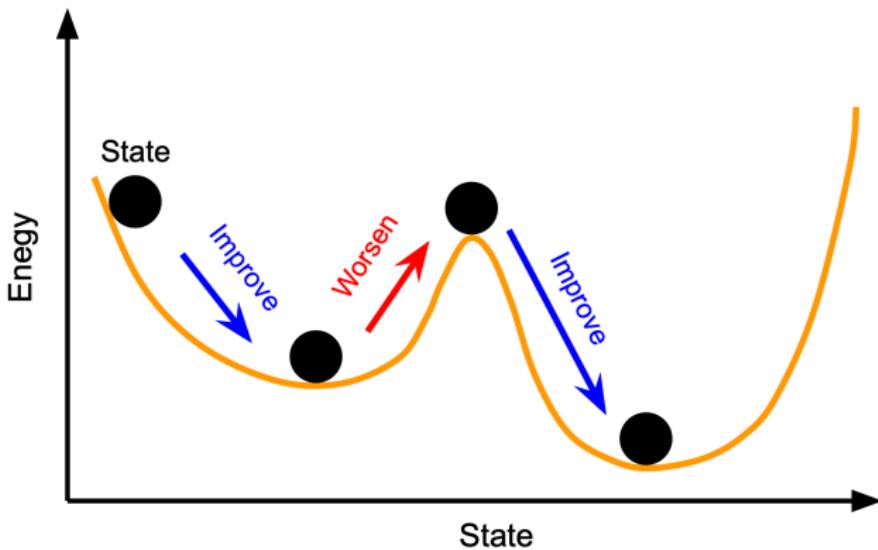


## MCTS Steps

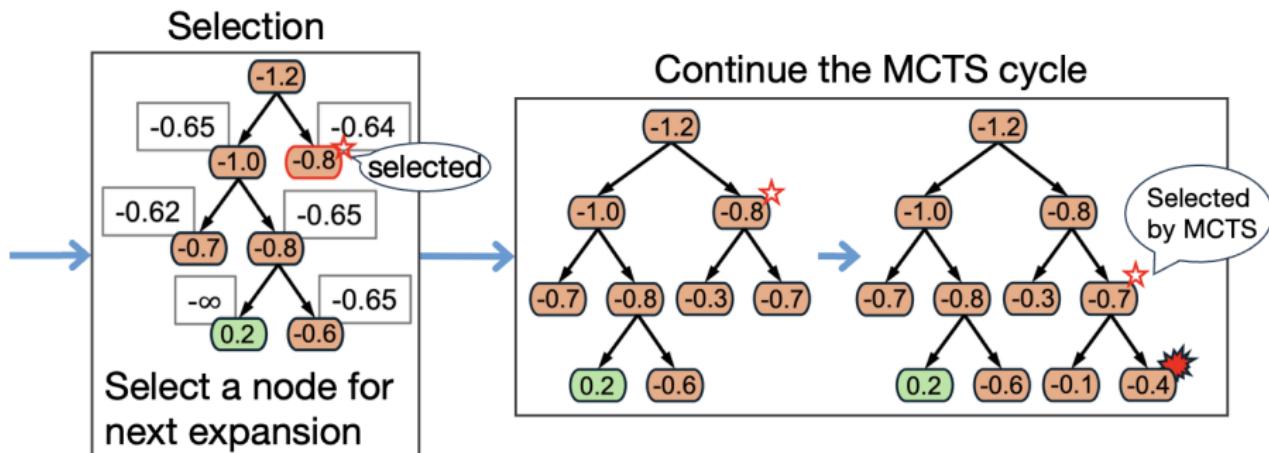
A MCTS cycle on BaB verification consists of 4 steps.



## Stochastic: *Oliva*<sup>SA</sup>



## MCTS Steps Cycle<sup>7</sup>



MCTS guides a smarter node expansion order based on CeP while speedup the counterexample finding.

<sup>7</sup> DATE25: Adaptive Branch-and-Bound Tree Exploration for Neural Network Verification: Kota Fukuda, **Guanqin Zhang**, Zhenya Zhang, Yulei Sui, Jianjun Zhao

## Evaluation

■ **Table 2** RQ1 – Overall comparison of different verification approaches, in terms of the number of solved problem instances and the time costs (in secs).

Model	BaB-baseline		$\alpha\beta$ -Crown		Neuralsat		Oliva <sup>GR</sup>		Oliva <sup>SA</sup>	
	Solved	Time	Solved	Time	Solved	Time	Solved	Time	Solved	Time
MNIST <sub>L2</sub>	96	126.41	87	51.32	99	32.37	95	96.79	99	57.76
MNIST <sub>L4</sub>	65	194.74	44	428.03	54	392.04	67	146.31	53	142.72
OVAL21 <sub>BASE</sub>	42	770.7	58	641.7	70	621.21	154	184.96	159	155.29
OVAL21 <sub>DEEP</sub>	33	694.74	55	552.24	55	539.59	92	250.72	87	261.68
OVAL21 <sub>WIDE</sub>	40	733.73	63	557.51	65	533.01	131	240.16	112	288.0

**Table 3** RQ1 – Pairwise comparison on the number of additional solved problem instances from all verification tasks. The number in each cell implies the number of problem instance solved by the approach of the row, but not solved by the approach of the column.

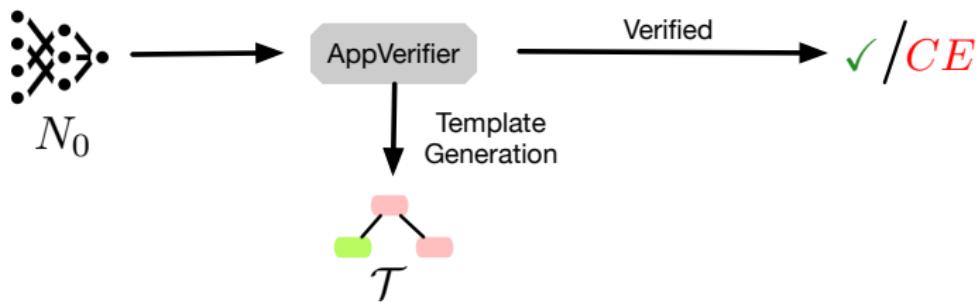
	BaB-Baseline	$\alpha\beta$ -Crown	NeuralSAT	Oliva <sup>GR</sup>	Oliva <sup>SA</sup>
BaB-Baseline	0	80	59	8	13
$\alpha\beta$ -Crown	111	0	11	23	39
NeuralSAT	126	47	0	30	46
Oliva <sup>GR</sup>	271	255	226	0	40
Oliva <sup>SA</sup>	247	242	213	11	0

# Incremental Neural Network Verification I

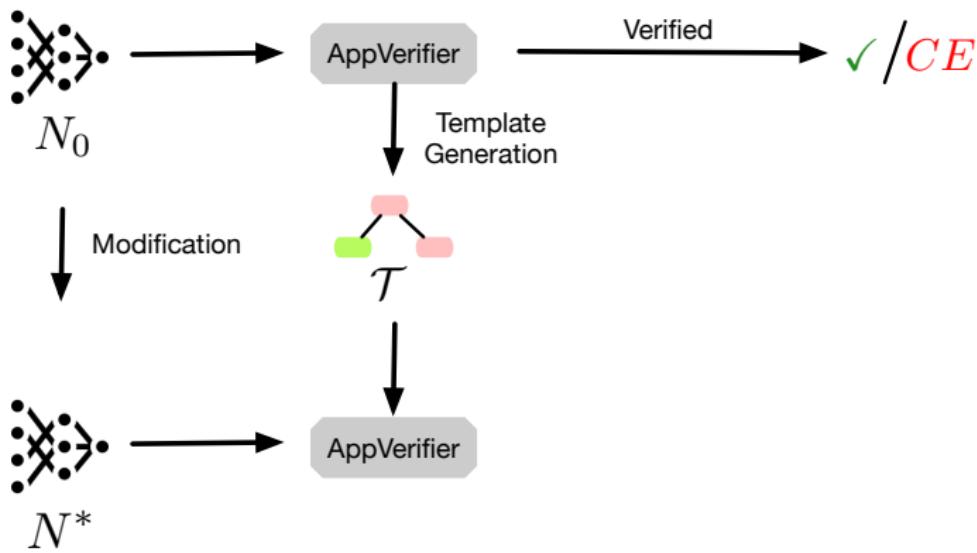
Existing Approach: IVAN



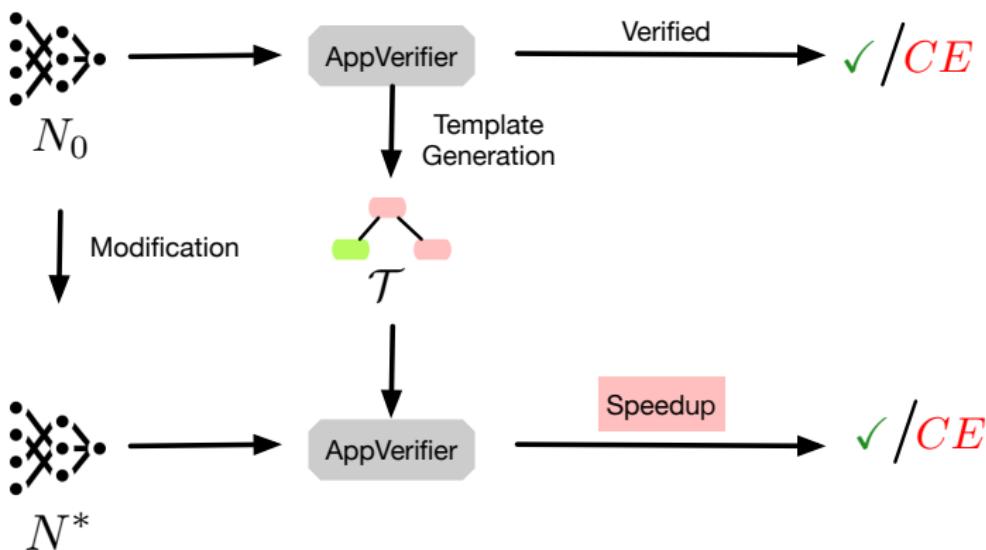
# Incremental Neural Network Verification II



# Incremental Neural Network Verification III

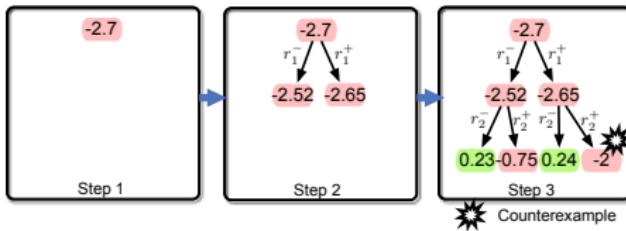
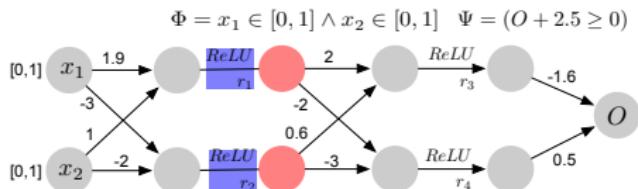


Incremental Neural Network Verification IV



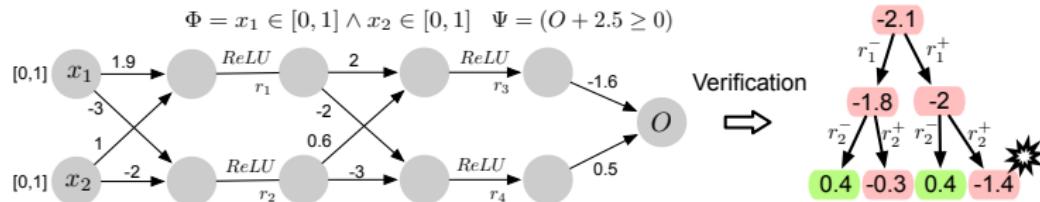
Noted:  $N^*$  has the identical structure to  $N$  but slightly differs in model parameters.

## Revist DNN Verification

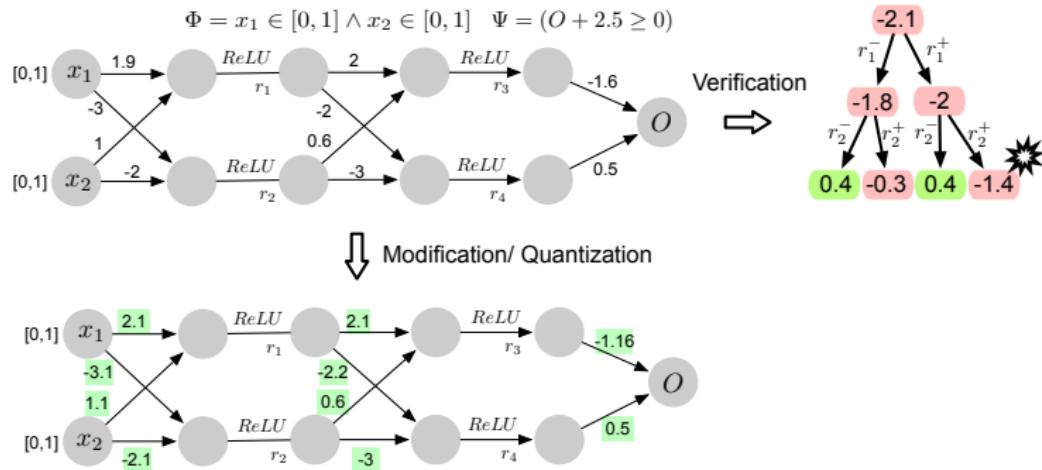


Specification Tree generated after Branch and Bound.

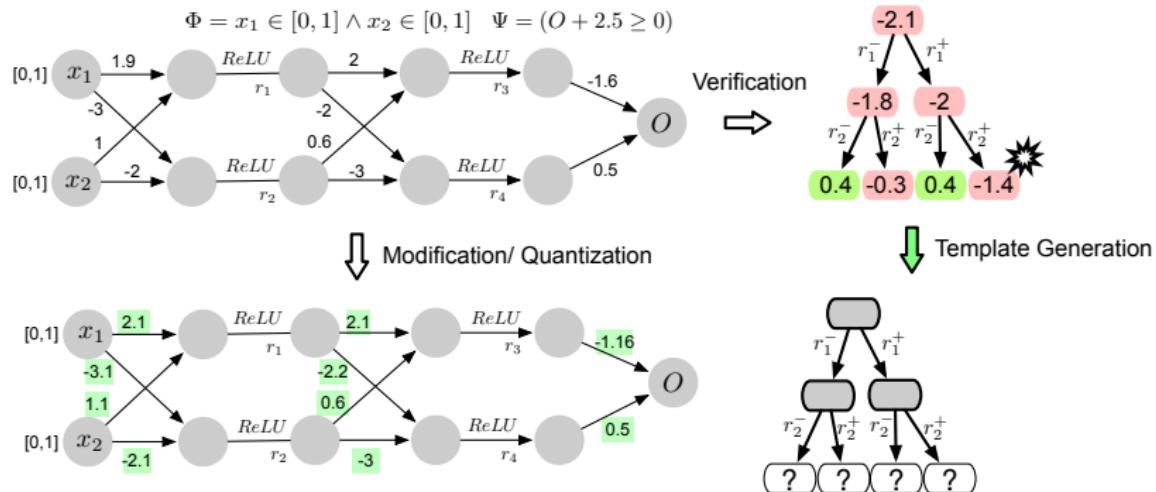
# Incremental Neural Network Verification



## Incremental Neural Network Verification

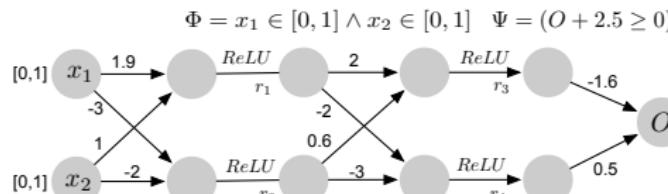


# Incremental Neural Network Verification

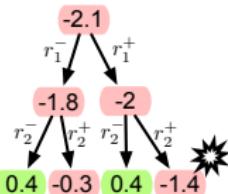


We only recompute the bounds in the **leaf nodes**.

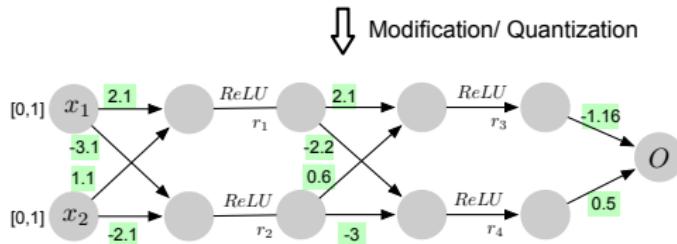
# Incremental Neural Network Verification



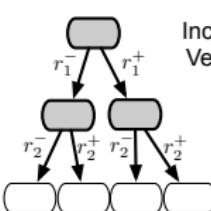
## Verification



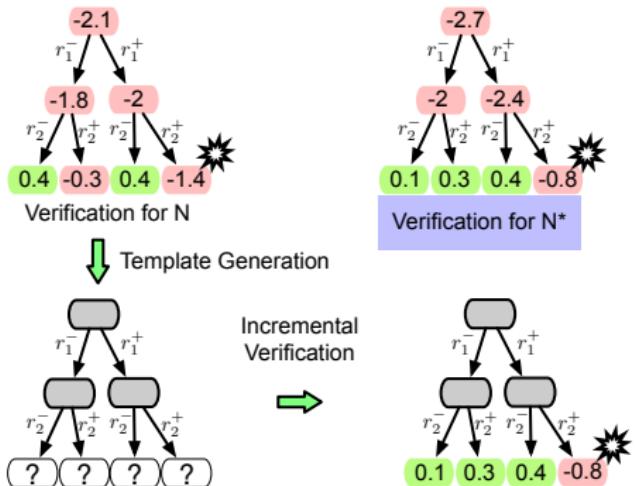
 Modification/ Quantization



Incremental  
Verification



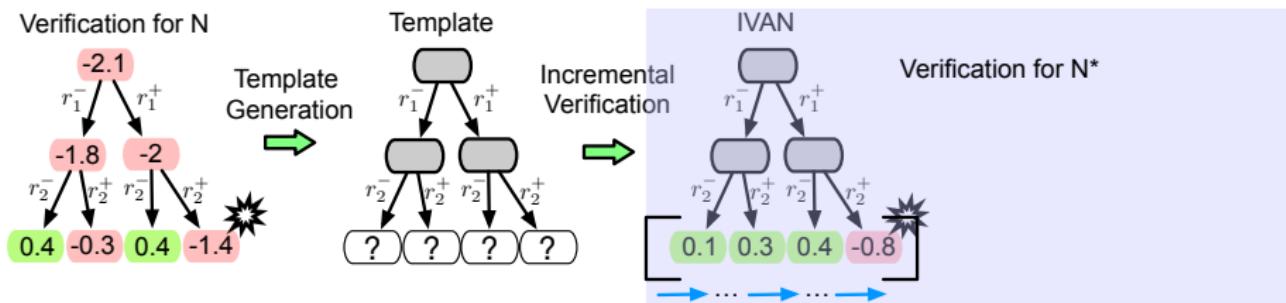
# Incremental Neural Network Verification (IVAN<sup>13.</sup>)



- IVAN builds on top of branch and bound based methods, and compares with non-incremental verification, gaining more scalability.
- IVAN saves extra **2** branching selection, **3** bounding computation.
- Assuming the non-incremental successfully verified  $2^{n+1} - 1$  nodes, the incremental verification only executes  $2^n$  bounding.

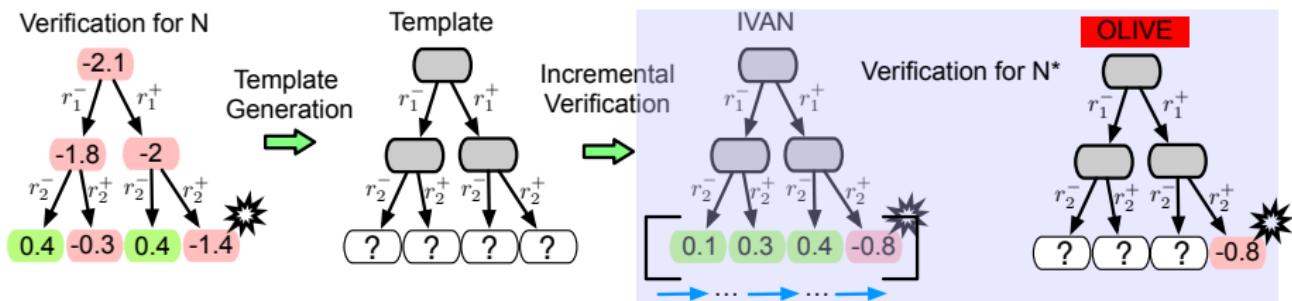
<sup>13</sup>Ugare, Shubham, et al. "Incremental verification of neural networks." Proceedings of the ACM on Programming Languages 7.PLDI (2023): 1920-1945.

## Aim of This Work: Order-leading Incremental Verification



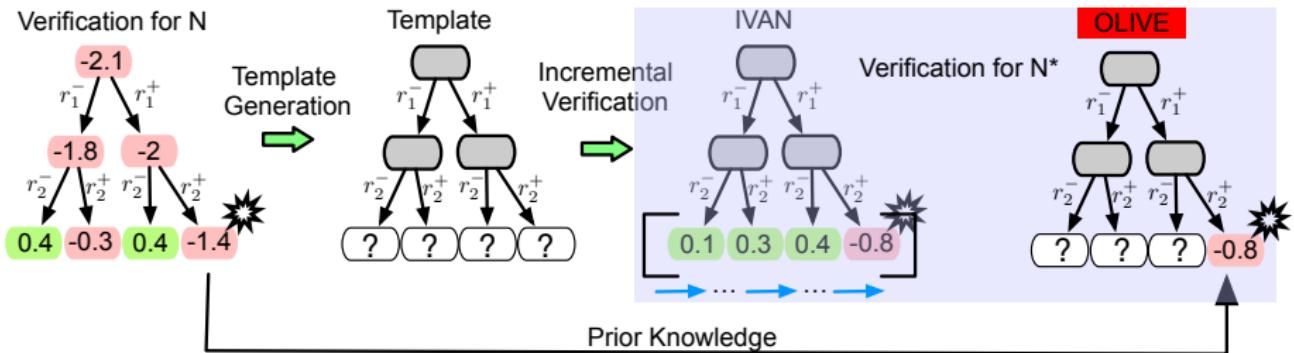
**Existing** work, such as IVAN, does not consider the **order** of the subproblem issues. It serves the template subproblem as in a **queue**, assuming as a “first come first serve”.

## Order-leading Incremental Verification (1): *Olive<sup>g</sup>*



We shortcuts the queue with less bounding times, in this case, we **speedup 4x**.

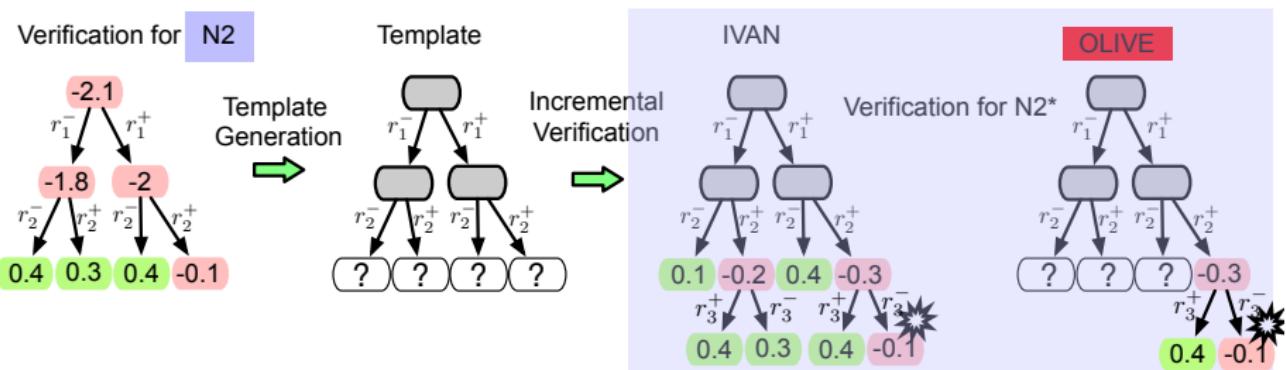
## Order-leading Incremental Verification (1): *Olive<sup>g</sup>*



We consider the **counterexample potentiality**<sup>14</sup> of each node in the specification tree.

<sup>14</sup>if a node  $\langle \Gamma_1, \hat{p}_1 \rangle$  is superior than another node  $\langle \Gamma_2, \hat{p}_2 \rangle$  in only one of the attributions (e.g.,  $|\Gamma_1| > |\Gamma_2|$  but  $\hat{p}_1 > \hat{p}_2$ ).

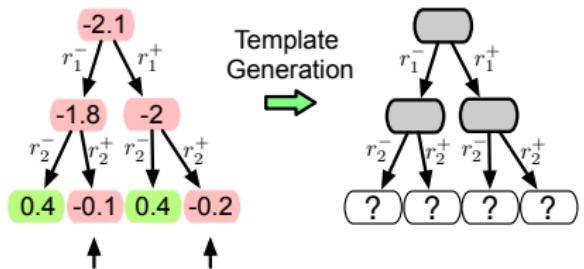
# Order-leading Incremental Verification (2): *Olive<sup>g</sup>*



This order is essential, because we exploit the most promising counterexample potential order for speedup. In this case, IVAN require **2** branching, and **8** bounding. *Olive<sup>g</sup>* only need **1** branching, and **3** bounding.

Olive<sup>g</sup> Insights

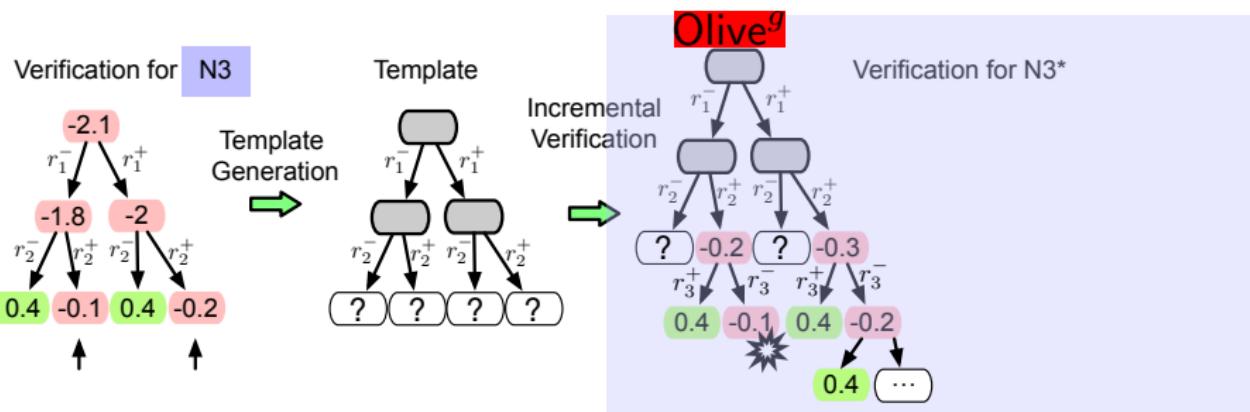
### Verification for N3



## Template

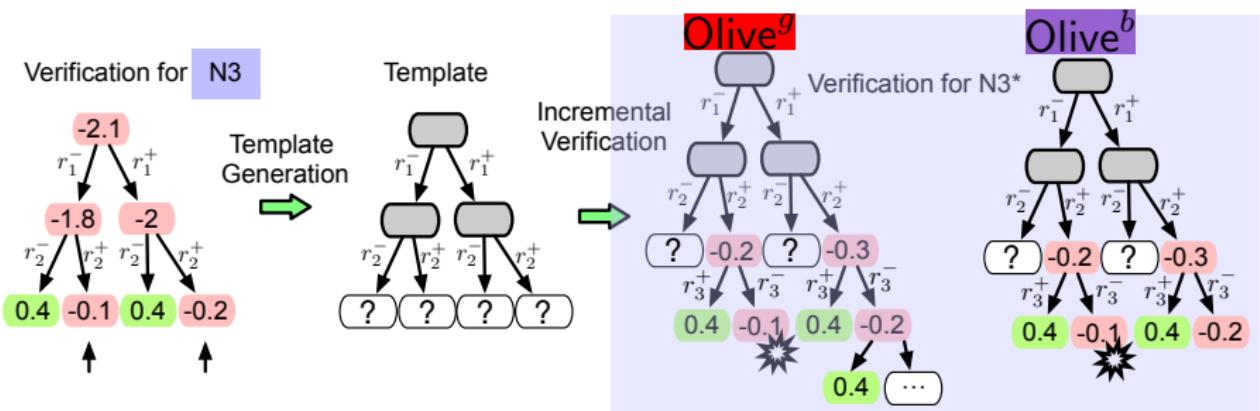


Olive<sup>g</sup> Insights



While the greedy strategy  $Olive^g$  favors exploitation of the suspicious sub-problem suggested by the specification tree  $\mathcal{T}_N$  of  $N$ , it can **fail** when the suggestion given by  $\mathcal{T}_N$  is not precise for  $N^*$ .

## Order-leading Incremental Verification (2): *Olive*<sup>b</sup>

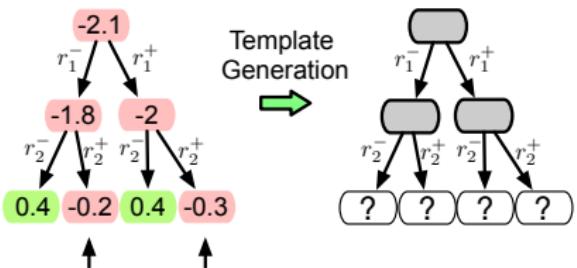


Namely, there could exist a different branch in  $N3^*$  other than the one being exploited, in which counterexamples are easier to be found. This raises the classic problem of “*exploration and exploitation*” trade-off in search-based techniques. In light of this, we propose a balanced strategy that models the incremental verification problem as a multi-armed bandit (MAB)<sup>15</sup>) problem.

<sup>15</sup>Slivkins, Aleksandrs. "Introduction to multi-armed bandits." Foundations and Trends in Machine Learning 12.1-2 (2019): 1-286.

# Olive<sup>b</sup> with MAB

Verification for N3

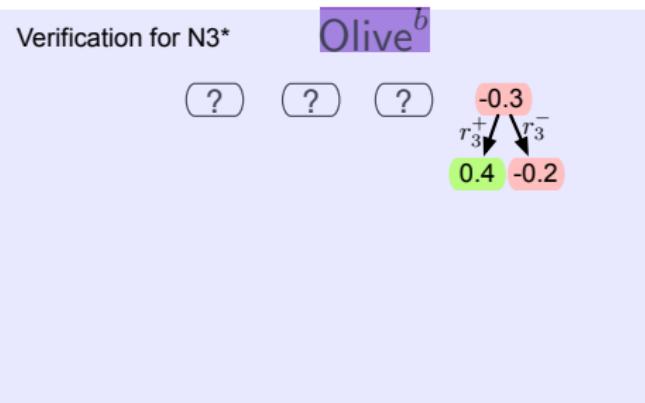
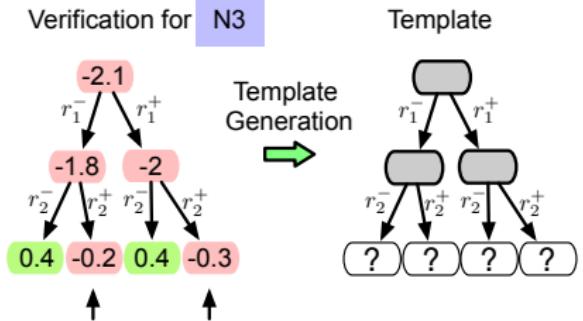


Verification for N3\*

Olive<sup>b</sup>

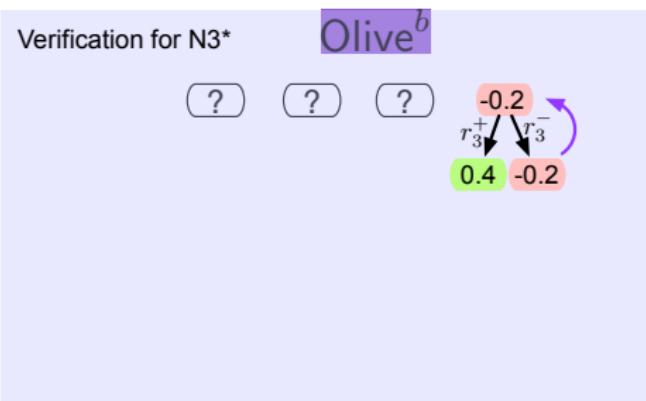
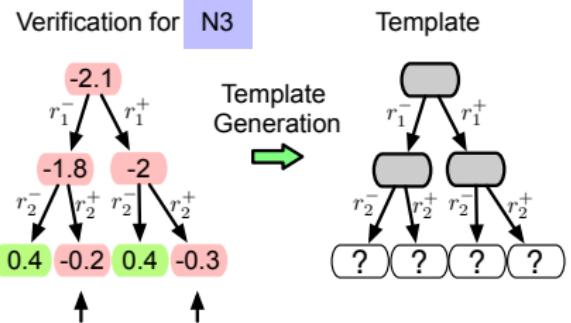
In Olive<sup>b</sup>, we consider the classic problem of "**exploration and exploitation**" trade-off in search-based techniques.

## Olive<sup>b</sup> with MAB - Compute Reward

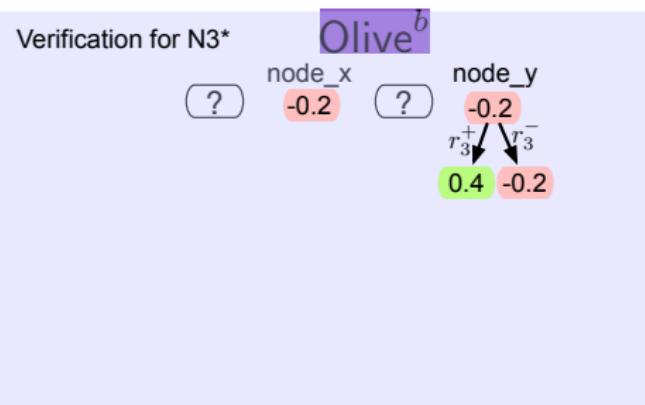
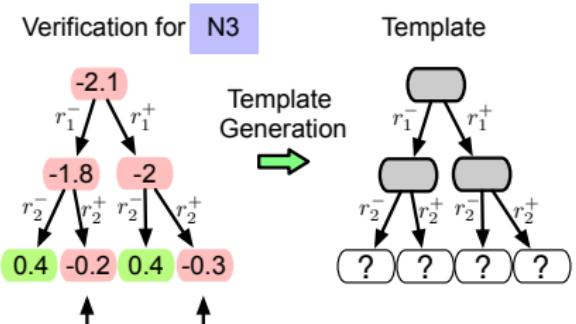


$$R(|\Gamma|, \hat{p}) := \begin{cases} +\infty & \text{if } \hat{p} < 0 \text{ and } Valid(\hat{x}, N^*, \Psi) \\ -\infty & \text{if } \hat{p} > 0 \\ \sigma \frac{|\Gamma|}{K} + (1 - \sigma) \frac{\hat{p}}{\hat{p}_{Min}} & \text{otherwise} \end{cases}$$

## *Olive<sup>b</sup>* with MAB - Back Propagation



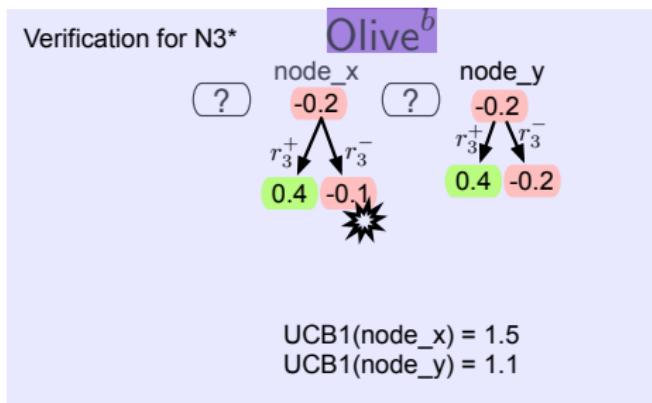
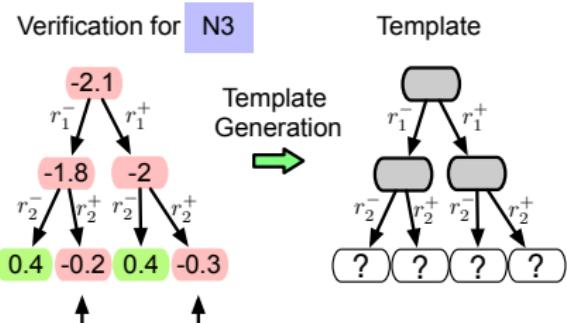
# Olive<sup>b</sup> with MAB - Arms Expansion



Conditional Expansion : if  $|W| \leq C \cdot (\sum_{\Gamma \in W} |T(\Gamma)|)^\alpha$  and  $|W| \neq |\mathcal{L}|$ <sup>16</sup>

<sup>16</sup>Here,  $W = 4$

# Olive<sup>b</sup> with MAB - UCB1



Select **next**:  $\Gamma^* \leftarrow \operatorname{argmax}_{\Gamma \in W} \left( R(\Gamma) + c \sqrt{\frac{2 \ln (\sum_{\Gamma \in W} |T(\Gamma)|)}{|T(\Gamma)|}} \right)$

# Evaluation

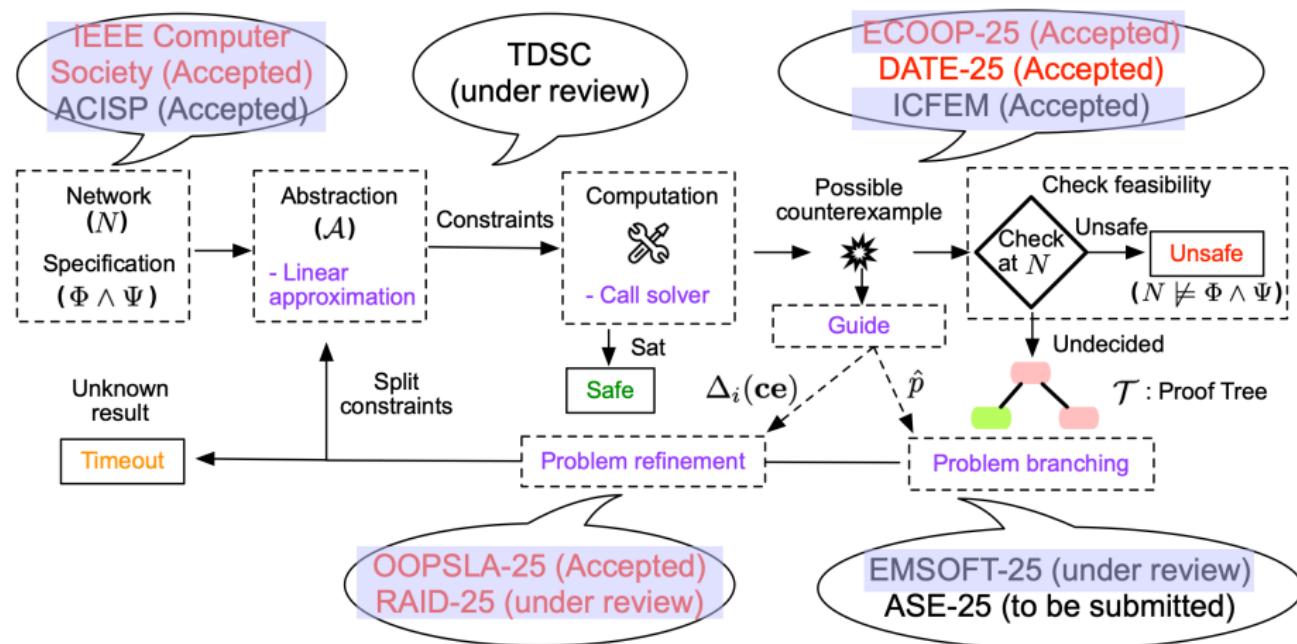
Model	Alteration	$\alpha\beta$ -Crown		Marabou		Ivan		Olive <sup>g</sup>		Olive <sup>b</sup>	
		Speedup	+Solved	Speedup	+Solved	Speedup	+Solved	Speedup	+Solved	Speedup	+Solved
MNIST <sub>L2</sub>	Quant(Int8)	1.08×	0	0.93×	0	3.59×	8	3.28×	5	2.51×	3
MNIST <sub>L2</sub>	Prune(7%)	1.05×	0	1.02×	0	1.60×	13	1.75×	10	1.73×	11
MNIST <sub>L2</sub>	Prune(12%)	2.28×	0	1.22×	0	2.31×	0	34.89×	1	35.98×	4
MNIST <sub>L4</sub>	Quant(Int8)	2.18×	0	1.28×	0	1.78×	4	2.15×	5	2.14×	6
MNIST <sub>L4</sub>	Prune(7%)	1.01×	0	0.99×	0	1.59×	4	2.16×	2	2.18×	3
MNIST <sub>L4</sub>	Prune(12%)	1.06×	0	1.1×	0	2.92×	2	14.74×	3	14.99×	23
OVAL21 <sub>BASE</sub>	Prune(0.1%)	1.01×	0	0.93×	0	1.74×	16	2.56×	11	2.58×	7
OVAL21 <sub>BASE</sub>	Prune(1%)	1.04×	0	0.99×	0	1.79×	17	2.37×	12	2.33×	7
OVAL21 <sub>BASE</sub>	Prune(3%)	1.00×	0	0.96×	0	1.77×	14	2.31×	10	2.29×	6
OVAL21 <sub>DEEP</sub>	Prune(0.1%)	1.02×	0	1.03×	0	1.91×	10	3.19×	8	3.21×	5
OVAL21 <sub>DEEP</sub>	Prune(1%)	1.56×	3	1.14×	0	1.92×	7	3.14×	5	3.18×	3
OVAL21 <sub>DEEP</sub>	Prune(3%)	1.03×	0	1.01×	0	1.88×	4	2.79×	2	2.78×	2
OVAL21 <sub>WIDE</sub>	Prune(0.1%)	1.0×	0	0.88×	0	1.98×	1	5.82×	1	5.98×	1
OVAL21 <sub>WIDE</sub>	Prune(1%)	0.98×	0	0.83×	0	1.74×	1	7.51×	1	7.34×	2
OVAL21 <sub>WIDE</sub>	Prune(3%)	1.02×	0	0.93×	0	1.88×	4	2.79×	2	2.78×	2

The performance comparison between three incremental verification approaches and the BaB baseline approach, in terms of the average speedup w.r.t. BaB over the solved verification problems and the number of the additional problems that are not solved by BaB but solved by other approaches.

# MileStone Summary

MileStones	Outcome
Complete major revision of the paper submitted to TOSEM	Resubmitted to TDSC
Complete major revision of the paper submitted to OOPSLA25	ACHIEVED
Complete current work (OLIVA) and submit it to CAV 2025	ACHIEVED in ECOOP
Participate in collaborators' work and submit it to VNN-COMP 2025	As planned and undergo
Draft thesis	As planned and undergo
Other Contributed works	Detailed in next slide

## Contributed works (All in One)



# Thesis Structure

## 1 Literature Review and Outline:

- Duration: March 2025 (1 month)
- Focus: Comprehensive literature review and creating a detailed outline for the thesis

## 2 Chapter 1 Refinement and Revision:

- Duration: April 2025 - May 2025 (2 months)
- IEEE work as in an introduction, refining and revising

## 3 Chapter 2 Major Work - 1:

- Duration: April 2025 - May 2025 (2 months)
- ECOOP, DATE

## 4 Chapter 3 Major Work - 2:

- Duration: April 2025 - May 2025 (2 months)
- OOPSLA

## 5 Chapter 4 Refinement and Revision:

- Duration: June 2025 (1 month)
- Focus: Refining and revising the all chapters of thesis

## 6 Full Thesis Initial Revision:

- Duration: June 2025 (1 month)
- Focus: First comprehensive revision of the entire thesis

# Questions?